

Deal Capture

Catching disputes on Wednesday, not Friday

Greenroom · Applied AI PM Case Study · Samadrita Roy

The slice, and what made me pick it

Before picking a slice, I queried the settlements table directly. Sixteen settlements are marked *status = 'disputed'*, and all sixteen have artist sign-off text that reads as positive resolution. “Looks good.” “ok wire monday.” “Sign off.” “OK.” A few thumbs up. None of them read as actual disputes.

That tells me the disputes are not actually happening at settlement. By the time the artist signs off, both sides have worked it out. The break point is upstream: the deal was never structurally captured at the front, so by show night the venue and the agent are operating on different facts. Fix capture and most of what currently shows up as a settlement-stage problem stops happening at the source. That is the leverage point, and that is the slice I built.

The flow lives at shows/[id]/capture-deal. Mariana pastes the agent’s deal email. An LLM extracts structured terms with source quotes, flags ambiguities, and attaches a dollar consequence to each so she can see what is at stake. Each flag resolves one of two ways: a pre-drafted clarifying email back to the agent, or an inline resolution if she already knows what was meant. Each resolution persists with a marker for which path was used. Once the deal is confirmed, the record is the source of truth, with both interpretations locked in writing and timestamped.

Design choices

The extraction prompt is the product. Most of the design work went into the prompt itself. It frames Mariana and the venue, names the deal types Greenroom actually sees (including the vs deal variants the current tool cannot handle), defines what counts as a money-changing ambiguity, and includes two grounded examples drawn from the data. The model surfaces uncertainty rather than hides it: low-confidence terms get flagged, and dollar estimates are left blank when the email does not give enough to compute one. The prompt lives in its own file so it iterates independently of the wiring.

Confirmation lives inside the record, not in a separate status field. Direct response to the database finding. Status fields drift the moment reality moves faster than the schema. So in this build, the agent’s confirmation is timestamped inside the deal_capture record itself. No separate flag, nothing to drift. The class of bug that produced the sixteen false disputes has no place to live in the new system.

Two resolution paths, one resting state. Forcing every flag through the agent loop would add friction Mariana does not need on the easy cases, and she would stop opening the tool. Letting her resolve inline when she already knows the answer keeps the surface useful daily, and a resolution_source field on each clarification records which path was used so reporting stays clean.

Trade-offs

Three deliberate sequencing choices, each made because the alternative was worse.

No settlement math in V1. Building a calculator on top of ambiguous inputs reproduces the situation Greenroom is in today. Trustworthy inputs first, deterministic math second. Math becomes a one-week build once capture lands.

No writeback to the legacy deals table. Pretending the old fields are now reliable re-introduces the exact drift this build exists to solve. The new record is the source of truth. Migrating reports off the legacy table is a clean, scoped follow-on, not a coupling that compromises the new schema.

Email plus one-click confirm instead of a full agent app. An agent-facing product would take weeks of account, onboarding, and permissions work that does not change the core idea. Email is where agents already live. The structured collaboration still happens, inside the deal_capture record both sides operate on.

What I cut, and the version of each I considered

A settlement calculator. The most visible pain in the brief, and the most tempting build. Cut because the math is the easy part. Mariana already has a calculator; the inputs are what is broken. Building a better UI on top of ambiguous inputs solves nothing.

Real-time deal prediction for Pri. Real PM value, and a natural V2. Cut because prediction without trusted inputs is a vitamin. Once capture is structured, prediction becomes a dashboard on real fields rather than a model trying to clean up someone else's data.

The 2 a.m. walkthrough UI. The emotional center of the brief. Cut because the walkthrough is hard only because the deal is ambiguous to begin with. Once capture is structured, the walkthrough becomes a reading of a trusted document, which is a smaller, downstream problem.

A dispute resolution UI. Disputes are the most expensive failure mode, so a dedicated surface looked compelling. Cut because disputes are a symptom. Build it first and you optimise the system for a state you do not want it to be in.

How I would validate this

Backtest first, then a scoped live test at The Crescent. Adoption is the bigger risk than accuracy, so the test is designed to read both.

Backtest. Run the extraction prompt across 24 months of notes_freetext entries and cross-reference flagged ambiguities against the historical dispute thread.

- Hypothesis: 70% or more of past disputes trace back to an ambiguity the system would have caught at capture.
- Kill criterion: below 50%. The wedge is weaker than claimed; rework before any live test.
- Watch for: flagged ambiguities that never became disputes. False-positive rate above 3x true-positive means the threshold is too loose and Mariana will tune it out.

Live test. One quarter at The Crescent, against the prior four quarters as baseline.

- **Primary:** end-of-show dispute rate. A drop of more than one third unlocks multi-venue rollout; no movement triggers a root-cause review for adoption versus design.
- **Adoption gate:** share of shows using the capture flow. Below 60% by week 4 means any movement in dispute rate is unreadable.
- **Leading indicator (week 2):** median time from email receipt to first sent clarification. Under 24h, the loop is working. Past 72h, Mariana is recording but not clarifying, and the two-sided thesis is failing in practice.

What I would ship next

- **Settlement math engine** that reads the confirmed deal_capture record and produces the artist payout, every line traceable to a clause in the structured deal and a receipt in the expense system. Inputs are now trustworthy, so math becomes deterministic.
- **Real outbound integration and reply parsing** to replace the mocked confirmation with the real loop through the venue's email provider.
- **Lightweight agent confirmation page** reached by a magic link in the clarification email, so agents confirm with one click instead of composing a reply.
- **Portfolio dashboard for Pri** showing, across the customer base, how many ambiguities the system flags at capture versus how many actually become disputes. Settlement craft becomes a number on a dashboard.